Metagenomics Workshop SciLifeLab Documentation
Release 1.0

Johannes Alneberg, John Larsson, Ino de Bruijn, Luisa Hugerth, A

December 11, 2014

Contents

This is a one day metagenomics workshop. We will discuss quality checking, assembly, taxonomic classification, binning and annotation of metagenomic samples.

Program:

- **Getting started**

- **Sessions**

Contents:

Introduction

A presentation for this workshop is available here.

Here is a link to the official homepage for the 2014 SciLifeLab workshop in metagenomics.

In this workshop we will be working with human associated metagenomes from the human microbiome project (HMP), following all the bioinformatic steps beginning with read sequence quality checking and ending with functional annotation of assembled contigs. You will be asked to choose one dataset to work with throughout the entire workshop and in the end we will compare results among different groups. The three datasets you can choose from is:

- Skin metagenome
- Tooth metagenome
- Gut metagenome

Connecting to UPPMAX

2.1 Connecting to UPPMAX

The first step of this lab is to open a ssh connection to UPPMAX. If you have a Mac, start the terminal (black screen icon). If you work on a PC, download and start MobaXterm (http://mobaxterm.mobatek.net). Now type (change username to your own username):

```
ssh -X username@milou.uppmax.uu.se
```

and give your password when prompted. As you type the password, nothing will show on screen. No stars, no dots. It is supposed to be that way. Just type the password and press enter, it will be fine. You should now get a welcoming message from Uppmax to show that you have successfully logged in.

2.2 Getting a node of your own

Usually you would do most of the work in this lab directly on one of the login nodes at uppmax, but we have arranged for you to have half of one node (=8 cores) each to avoid disturbances. To get this reservation you need to use the salloc command like this:

```
salloc -A g2014180 -t 08:00:00 -p core -n 8 --no-shell --reservation=g2014180_tue &
```

Now check which node you got (replace username with your uppmax user name) like this:

```
squeue -u username
```

The nodelist column gives you the name of the node that has been reserved for you (starts with "n"). Connect to that node using:

```
ssh -X nodename
```

Note: there is a uppmax specific tool called jobinfo that supplies the same kind of information as squeue that you can use as well ($ jobinfo -u username). You are now logged in to your reserved node, and there is no need for you to use the SLURM queuing system. You can now continue with the specific exercise instructions.

2.3 Load virtual environment

We have already installed all programs for you, all you have to do is load the virtual environment for this workshop. Once you are logged in to the server run:

```
source /proj/g2014180/metagenomics/virtenv/bin/activate
```

If you would have to, you deactivate the virtual environment with the command *deactivate*, but you don't have to do that yet.

NOTE: This is a python virtual environment. The binary folder of the virtual environment has symbolic links to all programs used in this workshop so you should be able to run those without problems.

2.4 Set sample variables

You will now have to make your decision on which kind of dataset you want to work with during this workshop. The choices you have are three different sampling sites on or within the human body:

- Gut
- Skin
- Teeth

Run only *one* of the following commands in the terminal

This will set the `SAMPLE` and `SAMPLE_ID` variables that will be used in the commands in the next steps of the tutorial. If for some reason you have to restart the terminal you will have to set these variable names again.

2.4.1 Gut

```
SAMPLE=gut
SAMPLE_ID=SRS011405
```

2.4.2 Teeth

```
SAMPLE=teeth
SAMPLE_ID=SRS014690
```

2.4.3 Skin

```
SAMPLE=skin
SAMPLE_ID=SRS015381
```

After you have chosen a sample you will create the file structure continuously throughout the workshop. This will make it possible for us to only use '$SAMPLE' in the commands, and it will automatically be changed to the sample type that you chose. Here you can see an overview of what this structure should look like at the end of the day (the "results" part of this structure).

Checking required software

An often occuring theme in bioinformatics is installing software. Here we wil go over some steps to help you check whether you actually have the right software installed. There's an optional excerise on how to install `sickle`.

3.1 Programs used in this workshop

The following programs are used in this workshop:

- Bowtie2
- Velvet
- samtools
- Picard
- Phylosift
- Fastqc
- Sortmerna
- Rdp_Classifier
- Krona
- Prokka
- MinPath
- BedTools

All programs are already installed, all you have to do is load the virtual environment for this workshop. Once you are logged in to the server run:

```
source /proj/g2014180/metagenomics/virtenv/bin/activate
```

If you would have to, you deactivate the virtual environment with the command *deactivate*, but you don't have to do that yet.

NOTE: This is a python virtual environment. The binary folder of the virtual environment has symbolic links to all programs used in this workshop so you should be able to run those without problems.

3.2 Using which to locate a program

An easy way to determine whether you have have a certain program installed is by typing:

```
which programname
```

where `programname` is the name of the program you want to use. The program `which` searches all directories in `$PATH` for the executable file `programname` and returns the path of the first found hit. This is exactly what happens when you would just type `programname` on the command line, but then `programname` is also executed. To see what your `$PATH` looks like, simply `echo` it:

```
echo $PATH
```

For more information on the `$PATH` variable see this link: http://www.linfo.org/path_env_var.html.

3.3 Check all programs in one go with which

To check whether you have all programs installed in one go, you can use `which`. In order to do so we will iterate over all the programs and call `which` on each of them. First make a variable containing all programs separated by whitespace:

```
$ req_progs="bowtie2 bowtie2-build velveth velvetg parallel samtools interleave-reads.py phylosift fa
$ echo $req_progs
bowtie2 bowtie2-build velveth velvetg parallel samtools interleave-reads.py phylosift fastqc sortmern
```

Now iterate over the variable `req_progs` and call which:

```
$ for p in $req_progs; do which $p || echo $p not in PATH; done
/proj/g2014180/metagenomics/virtenv/bin/bowtie2
/proj/g2014180/metagenomics/virtenv/bin/bowtie2-build
/proj/g2014180/metagenomics/virtenv/bin/velveth
/proj/g2014180/metagenomics/virtenv/bin/velvetg
/sw/parallel/gnuparallel/20140222/bin/parallel
/proj/g2014180/metagenomics/virtenv/bin/samtools
/proj/g2014180/metagenomics/virtenv/bin/interleave-reads.py
/proj/g2014180/metagenomics/phylosift_v1.0.1/phylosift
/sw/apps/bioinfo/fastqc/0.11.2/milou/fastqc
/proj/g2014180/metagenomics/virtenv/bin/sortmerna
/proj/g2014180/metagenomics/virtenv/bin/prokka
/proj/g2014180/metagenomics/virtenv/bin/MinPath1.2.py
/proj/g2014180/metagenomics/virtenv/bin/bedtools
```

In Unix-like systems a program that sucessfully completes it tasks should return a zero exit status. For the program `which` that is the case if the program is found. The `||` character does not mean *pipe the output onward* as you are probably familiar with (otherwise see http://tldp.org/HOWTO/Bash-Prog-Intro-HOWTO-4.html), but checks whether the program before it exists succesfully and executes the part behind it if not.

If any of the installed programs is missing, try to install them yourself or ask. If you are having troubles following these examples, try to find some bash tutorials online next time you have some time to kill. Educating yourself on how to use the command line effectively increases your productivity immensely.

Some bash resources:

- Excellent bash tutorial http://tldp.org/HOWTO/Bash-Prog-Intro-HOWTO.html

- Blog post on pipes for NGS http://www.vincebuffalo.com/2013/08/08/the-mighty-named-pipe.html

- Using bash and GNU parallel for NGS http://bit.ly/gwbash

3.4 (Optional excercise) Install sickle by yourself

Follow these steps only if you want to install `sickle` by yourself. Installation procedures of research software often follow the same pattern. Download the code, *compile* it and copy the binary to a location in your `$PATH`. The code for sickle is on https://github.com/najoshi/sickle. I prefer *compiling* my programs in `~/src` and then copying the resulting program to my `~/bin` directory, which is in my `$PATH`. This should get you a long way:

```
mkdir -p ~/src

# Go to the source directory and clone the sickle repository
cd ~/src
git clone https://github.com/najoshi/sickle
cd sickle

# Compile the program
make

# Create a bin directory
mkdir -p ~/bin
cp sickle ~/bin
```

Quality Control

The first step of any sequencing project is to do quality control of your reads and remove (trim) low quality bases from the end of the read. In this exercise, you will work with Illumina data from the Human Microbiome Project that has already been trimmed. We still want to check the quality of reads, though.

In this part of the metagenomics workshop we will learn how to:

- Check the quality of your raw sequencing data
- Perform quality trimming using sickle

The workshop has the following exercises:

4.1 Quality Control with fastqc

In this excercise you will use fastqc to investigate the quality of your sequences using a nice graphical summary output.

4.1.1 Retrieving your data

For the first step, make a workshop folder in your home directory and move into it:

```
mkdir -p ~/mg-workshop
cd ~/mg-workshop
```

Inside it, make a folder for your input files:

```
mkdir -p ~/mg-workshop/data/$SAMPLE/reads/1M/
cd ~/mg-workshop/data/$SAMPLE/
```

Now make a copy of the files you want to work on: gut, skin or teeth datasets. These files were originally taken from the Human Microbiome Project and then subsampled to include only 1 million reads each. You can copy these files from the project directory:

```
cp /proj/g2014180/nobackup/metagenomics-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.1.fastq ~/mg-w
cp /proj/g2014180/nobackup/metagenomics-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.2.fastq ~/mg-w
```

You will now have two files in your reads directory: one for the forward reads `*_1.fastq` and one for the reverse reads `*_2.fastq`.

4.1.2 Fastqc

We will now use FastQC to generate a report about the quality of our sequencing reads. For most programs and scripts in this workshop, you can see their instructions by typing their name in the terminal followed by the flag -h. There are many options available, and we'll show you only a few of those.

First, make a folder to keep your quality control results:

```
mkdir -p ~/mg-workshop/results/quality_check/$SAMPLE/
```

Now, run fastqc for each file:

```
fastqc -o ~/mg-workshop/results/quality_check/$SAMPLE/ --nogroup reads/1M/${SAMPLE_ID}_1M.1.fastq rea
```

FastQC will generate two files for each input file, one compressed, and one not. To view your files, copy the html results into your local computer and open them with a browser.

From **your own shell (not inside Uppmax - open a new terminal window)**:

```
mkdir -p ~/mg-workshop/
cd ~/mg-workshop/
scp -r username@milou.uppmax.uu.se:~/mg-workshop/results/quality_check/*/*html .
```

Instead of username, type your own username!

Now open the reports. Make sure you understand the results. Do they look ok? Is there a difference between forward and reverse? Are there any warnings or errors? What do they mean? Do you have adapter sequences in your reads?

4.2 Optional: Quality trimming Illumina paired-end reads

In this excercise you will learn how to quality trim Illumina paired-end reads. The most common Next Generation Sequencing (NGS) technology for metagenomics. The reads from the HMP are already quality trimmed. However, if you have time and want to try it out for yourself, you can run some more stringent quality trimming on them and see what happens.

4.2.1 Running sickle on a paired end library

For quality trimming Illumina paired end reads we use the library sickle which trims reads from 3' end to 5' end using a sliding window. If the mean quality drops below a specified number the remaining part of the read will be trimmed.

As a default, sickle trims a read at the point needed to maintain its average quality over 20. It also discards reads that are shorter than 20 bp. These are very good default values, but in this extra exercice you're welcome to change the values of these parameters using the -q and -l flags.

You can use the same qc directory as before for this step, since these reads won't be further processed.

Run sickle:

```
mkdir -p ~/mg-workshop/results/quality_check/sickle/$SAMPLE
sickle pe \
        -f ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.1.fastq \
        -r ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.2.fastq \
        -t sanger \
        -o ~/mg-workshop/results/quality_check/sickle/$SAMPLE/qtrim.1.fastq \
        -p ~/mg-workshop/results/quality_check/sickle/$SAMPLE/qtrim.2.fastq \
        -s ~/mg-workshop/results/quality_check/sickle/$SAMPLE/qtrim.unpaired.fastq \
        -q 20 -l 20
```

Chek what files have been generated. Do you understand each of them?

Question: How many paired reads are left after trimming? How many singletons?

Question: What are the different quality scores that sickle can handle? Why do we specify -t sanger here?

4.2.2 Run fastqc again

We would like to see if sickle has done a good job, we do so by asserting the quality of the reads again with fastqc. Please refer to the FastQC exercise for instructions on how to do this.

Question: Does the quality improve much?

4.2.3 Trimming adapter sequence

To remove adapter sequences from your reads one can use cutadapt. We won't do that in this workshop.

At least a basic knowledge of how to work with the command line is required otherwise it will be very difficult to follow some of the examples. Have fun!

Community analysis using 16S reads

In this part of the metagenomics workshop we will learn how to:

The workshop has the following exercise:

5.1 Community analysis using rRNA gene reads

In this exercise we will analyse the taxonomic composition of your sample by utilising 16S rRNA encoding sequence reads. Partial 16S rRNA genes will be extracted from the reads using the program sortmeRNA and these will subsequenctly be classified using the RDP classifier. Finally, the results will be visualised with the interactive program KRONA.

5.1.1 SortmeRNA

We will extract 16S rRNA encoding reads using sortmeRNA which is one of the fastest software for this. We start by making the necessary folders and assigning all necessary databases to a variable called DB:

```
mkdir -p ~/mg-workshop/results/phylogeny/16S/$SAMPLE
cd ~/mg-workshop/results/phylogeny/16S/$SAMPLE
ln -s ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.1.fastq reads.1.fastq
ln -s ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.2.fastq reads.2.fastq
DB=/proj/g2013206/metagenomics/src/sortmerna-1.9/rRNA_databases/silva-arc-16s-database-id95.fasta,/p
```

SortMeRNA has built-in multithreading support that we will use for parallelization (-a). We still have to launch one sample at a time, though:

```
for readfile in reads.*.fastq;
do sortmerna --reads $readfile --ref $DB --fastx --aligned ${readfile}_rrna -v -a 2;
done
```

5.1.2 RDP classifier

sortmeRNA outputs the reads, or part of reads, that encode rRNA in a fasta file. These rRNA sequences can be classified in many ways. One option is blasting them against a suitable database. Here we use a simple and fast method, the classifier tool at RDP (the Ribosomal Database Project). This uses a naïve bayesian classifier trained on kmer frequencies of many sequences of defined taxonomies. It gives bootstrap support values for each taxonomic level - usually, the support gets lower the further down the hierarchy you go. Genus level is the lowest level provided. You can use the web service if you prefer, and upload each file individually, or you can use the uppmax installation of RDP classifier like this:

```
for file in *_rrna*.fastq;
do name=$(basename $file);
java -Xmx1g -jar /proj/g2014180/metagenomics/virtenv/rdp_classifier_2.6/dist/classifier.jar classify
done
```

5.1.3 Krona

To get a graphical representation of the taxonomic classifications you can use Krona, which is an excellent program for exploring data with hierarchical structures in general. The output file is an html file that can be viewed in a browser. Again make a directory for Krona

And run Krona, specifycing the name of the output file (-o), the minimum bootstrap support to use (-m) and that the two input files should be treated as only one (-c):

```
ktImportRDP -o 16S.tax.html -m 50 -c reads.1.fastq_rrna.fastq.class.tsv  reads.2.fastq_rrna.fastq.cla
```

Copy the resulting file 16S.tax.html to your local computer with scp and open it a browser, like you did for the FastQC output.

Question: What's the dominant type of organisms found in your sample?

Have fun!

Metagenomic Assembly

In this part of the metagenomics workshop we will learn how to:

- Perform assemblies with velvet

The part has the following exercise:

6.1 Assembling reads with Velvet

In this exercise you will learn how to perform an assembly with Velvet. Velvet takes your reads as input and turns them into contigs. It consists of two steps. In the first step, `velveth`, the de Bruijn graph is created. Afterwards the graph is traversed and contigs are created with `velvetg`. When constructing the de Bruijn graph, a *kmer* has to be specified. Reads are cut up into pieces of length *k*, each representing a node in the graph, edges represent an overlap (some de Bruijn graph assemblers do this differently, but the idea is the same). The advantage of using kmer overlap instead of read overlap is that the computational requirements grow with the number of unique kmers instead of unique reads. A more detailed explanation can be found in this paper.

You can test different kmer lengths, as long as they're odd numbers. A good margin is to have the kmer length between 21 and 51. We'll then look at a few statistics on the assembly; if you're choice of kmer wasn't good, you might have to run another assembly (but this is very fast).

6.1.1 Pick your kmer

Fill in which kmer you want to do in the Google doc. Later we will compare the results from the different kmers for each group.

6.1.2 velveth

Create the graph data structure with `velveth`. First create a directory with symbolic links to the pairs that you want to use:

```
mkdir -p ~/mg-workshop/results/assembly/$SAMPLE/
cd ~/mg-workshop/results/assembly/$SAMPLE/
ln -s ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.1.fastq pair1.fastq
ln -s ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.2.fastq pair2.fastq
```

Create a directory for the kmer of your choice. **Replace N with the kmer length below**:

```
mkdir ${SAMPLE}_N
```

The reads need to be interleaved for `velveth`:

```
interleave-reads.py -o pair.fastq pair1.fastq pair2.fastq
```

Run velveth, **replacing N with the kmer length you chose**:

```
velveth ${SAMPLE}_N N -fastq -shortPaired pair.fastq
```

Check what directories have been created:

```
ls
```

6.1.3 velvetg

To get the actual contigs you will have to run `velvetg` on the created graph. You can vary options such expected coverage and the coverage cut-off if you want, but we do not do that in this tutorial. We only choose not to do scaffolding. Again **replace N for your current kmer length**:

```
velvetg ${SAMPLE}_N -scaffolding no
```

6.1.4 assemstats

After the assembly one wants to look at the length distributions of the resulting assemblies. You can use the `assemstats` script for that:

```
assemstats 100 ${SAMPLE}_N/contigs.fa
```

Try to find out what each of the stats represent by varying the cut-off. One of the most often used statistics in assembly length distribution comparisons is the *N50 length*, a weighted median of the length, where you weigh each contig by its length. This way, you assign more weight to larger contigs. Fifty per cent of all the bases in the assembly are contained in contigs shorter or equal to N50 length. Add your results to the Google doc.

Question: What are the important length statistics? Do we prefer sum over length? Should it be a combination?

6.1.5 (Optional) Ray

Try to use the Ray assembler instead. Ray was made to play well with metagenomics. Furthermore it uses MPI to distribute the computation over multiple computational nodes and/or cores. You can run Ray on 16 cores with the command:

```
mkdir -p ~/mg-workshop/results/assembly/ray/$SAMPLE/
module unload intel
module load gcc openmpi/1.7.5
rm -rf ~/mg-workshop/results/assembly/ray/$SAMPLE/${SAMPLE}_N
time mpiexec -n 16 Ray -k N -p ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.{1,2}.fastq \
    -o ~/mg-workshop/results/assembly/ray/$SAMPLE/${SAMPLE}_N
module unload gcc
module load intel
```

Replace N again with your chosen kmer. There is another sheet where you can add the Ray assembly results.

Question: How do Ray's results compare to those from Velvet?

Have fun!

Taxonomic Classification

- Taxonomic annotation of contigs using Phylosift

Contents:

7.1 Phylogenetic Classification using Phylosift

In this section investigate our contigs with Phylosift to see which species they originate from.

7.1.1 Phylosift

Phylosift is software created for the purpose of determining the phylogenetic composition of your metagenomic data. It uses a defined set of genes to predict the taxonomy of each sequence in your dataset. You can read more about how this works here: http://phylosift.wordpress.com. Lets prepare for the phylosift run:

```
mkdir -p ~/mg-workshop/results/phylogeny/phylosift/$SAMPLE
cd ~/mg-workshop/results/phylogeny/phylosift/$SAMPLE
ln -s /proj/g2014180/nobackup/metagenomics-workshop/results/assembly/$SAMPLE/${SAMPLE}_31/contigs.fa
```

You can run phylosift using the following command:

```
cd ~/mg-workshop/results/phylogeny/phylosift/$SAMPLE
phylosift all -f --output phylosift_output contigs.fa &
```

You can check the progress of the phylosift run by running the following command (You can repeat this):

```
cat ~/mg-workshop/results/phylogeny/phylosift/$SAMPLE/phylosift_output/run_info.txt
```

Unfortunately, the phylosift run is taking a long time. So if you don't have time to wait for it, in order to have some nice results to study, we'll copy the results from the project directory:

```
cp -r /proj/g2014180/nobackup/metagenomics-workshop/results/phylogeny/phylosift/$SAMPLE/phylosift_out
```

When the phylosift run is completed, browse the output directory:

```
ls ~/mg-workshop/results/phylogeny/phylosift/$SAMPLE/phylosift_output/
```

All of these files are interesting, but the most fun one is the html file, so lets download this to your own computer and have a look. **Again, switch to a terminal where you're not logged in to UPPMAX**:

```
mkdir ~/mg-workshop/
scp username@milou.uppmax.uu.se:~/mg-workshop/results/phylogeny/phylosift/phylosift_output/*.html ~/m
```

Question: Compare with the taxonomic results from the 16S analysis. Do the results match? If not, what could be the explanation for the differences?

Functional Annotation

This part of the workshop has the following exercises:

1. Gene annotation pipeline - PROKKA

2. Predict metabolic pathways using MinPath

3. Quantify genes by mapping reads to the assembly

4. Explore gene annotation using KRONA

Contents:

8.1 Annotating the assembly using the PROKKA pipeline

Now that you have assembled the data into contigs the next natural step to do is annotation of the data, i.e. finding the genes and doing functional annotation of those. A range of programs are available for these tasks but here we will use PROKKA, which is essentially a pipeline comprising several open source bioinformatic tools and databases. PROKKA automates the process of locating open reading frames (ORFs) and RNA regions on contigs, translating ORFs to protein sequences, searching for protein homologs and producing standard output files. For gene finding and translation, PROKKA makes use of the program Prodigal. Homology searching (BLAST, hmmscan) is then performed using the translated protein sequences as queries against a set of public databases (CDD, PFAM, TIGRFAM) as well as custom databases that come with PROKKA.

Set up the necessary files and run PROKKA:

```
mkdir -p ~/mg-workshop/results/functional_annotation/prokka/
cd ~/mg-workshop/results/functional_annotation/prokka/
ln -s /proj/g2014180/nobackup/metagenomics-workshop/results/assembly/$SAMPLE/${SAMPLE}_31/contigs.fa
prokka contigs.fa --outdir $SAMPLE --norrna --notrna --metagenome --cpus 16
cd $SAMPLE
```

PROKKA produces several types of output, such as:

- the **GFF** file, which is a standardised, tab delimited, file type for showing annotations

- the Genbank (**GBK**) file, which is a more detailed description of nucleotide sequences and the genes encoded on these.

An explanation of the gff format can be found at http://genome.ucsc.edu/FAQ/FAQformat.html.

An explanation of the Genbank format can be found at http://www.ncbi.nlm.nih.gov/Sitemap/samplerecord.html

When your dataset has been annotated you can view the annotations directly in the GFF file by doing:

```
less -S PROKKA_11252014.gff
```

Question: How many coding regions were found by Prodigal? Hint: use grep -c

Some genes in your dataset should now contain annotations from several databases, for instance enzyme and COG (Clusters of Orthologous Groups) identifiers.

Question: How many of the coding regions were given an enzyme identifier? How many were given a COG identifier?

In the downstream analyses we will quantify and compare the abundance of enzymes and metabolic pathways, as well as COGs in the different samples. To do this, we will first extract lists of the genes with enzyme and COG IDs from the GFF file that was produced by PROKKA. First we extract enzyme numbers for genes using pattern matching:

```
grep "eC_number=" PROKKA_11252014.gff | cut -f9 | cut -f1,2 -d ';'| sed 's/ID=//g'| sed 's/;eC_number
```

Then we do the same for COG identifiers:

```
egrep "COG[0-9]{4}" PROKKA_11252014.gff | cut -f9 | sed 's/.\+COG\([0-9]\+\);locus_tag=\(PROKKA_[0-9]
```

The COG table we will save for later. Next up is to predict pathways in the sample based on the enzymes annotated by PROKKA.

8.2 Predicting metabolic pathways using MinPath

Metabolic pathways are made up of enzymes that catalyze various reactions. Depending on how pathways are defined, they may contain any number of enzymes. A single enzyme may also be part of one or several pathways. One way of predicting metabolic pathways in a sample is to simply consider all the pathways that a set of enzymes are involved in. This may however overestimate pathways, for instance if only a few of the enzymes required for a pathway are annotated in the sample.

Here we will predict pathways using the program MinPath to get conservative estimate of the pathways present. MinPath only considers the minimum number of pathways required to explain the set of enzymes in the sample. As input, MinPath requires 1) a file with gene identifiers and enzyme numbers, separated by tabs, and 2) a file that links each enzyme to one or several pathways. The first of these we produced above using pattern matching from the PROKKA gff file. The second file exist in two versions, one that links enzymes to pathways defined in the Metacyc database and one that links enzymes to pathways defined in the KEGG database.

First we make sure that all the required files are available:

```
mkdir -p ~/mg-workshop/results/functional_annotation/minpath/$SAMPLE/
cd ~/mg-workshop/results/functional_annotation/minpath/$SAMPLE/
mkdir -p ~/mg-workshop/reference_db/
cp -r /proj/g2014180/nobackup/metagenomics-workshop/reference_db/cog ~/mg-workshop/reference_db/
cp -r /proj/g2014180/nobackup/metagenomics-workshop/reference_db/kegg ~/mg-workshop/reference_db/
cp -r /proj/g2014180/nobackup/metagenomics-workshop/reference_db/metacyc ~/mg-workshop/reference_db/
ln -s ~/mg-workshop/results/functional_annotation/prokka/$SAMPLE/PROKKA.$SAMPLE.ec
```

Run MinPath with this command to predict Metacyc pathways:

```
MinPath1.2.py -any PROKKA.$SAMPLE.ec -map ~/mg-workshop/reference_db/metacyc/ec.to.pwy -report PROKKA
```

And to predict KEGG pathways:

```
MinPath1.2.py -any PROKKA.$SAMPLE.ec -map ~/mg-workshop/reference_db/kegg/ec.to.pwy -report PROKKA.$S
```

Take a look at the report files:

```
less -S PROKKA.$SAMPLE.metacyc.minpath
```

Question: How many Metacyc and KEGG pathways did MinPath predict in your sample? How many were predicted if you had counted all possible pathways as being present? (HINT: look for the 'naive' and 'minpath' tags)

8.3 Mapping reads and quantifying genes

8.3.1 Overview

So far we have only got the number of genes and annotations in the sample. Because these annotations are predicted from assembled reads we have lost the quantitatve information for the annotations. So to actually quantify the genes we will map the input reads back to the assembly.

There are many different mappers available to map your reads back to the assemblies. Usually they result in a SAM or BAM file (http://genome.sph.umich.edu/wiki/SAM). Those are formats that contain the alignment information, where BAM is the binary version of the plain text SAM format. In this tutorial we will be using bowtie2 (http://bowtie-bio.sourceforge.net/bowtie2/index.shtml).

The SAM/BAM file can afterwards be processed with Picard (http://picard.sourceforge.net/) to remove duplicate reads. Those are likely to be reads that come from a PCR duplicate (http://www.biostars.org/p/15818/).

BEDTools (http://code.google.com/p/bedtools/) can then be used to retrieve coverage statistics.

There is a script available that does it all at once. Read it and try to understand what happens in each step:

```
less `which map-bowtie2-markduplicates.sh`
map-bowtie2-markduplicates.sh -h
```

Bowtie2 has some nice documentation: http://bowtie-bio.sourceforge.net/bowtie2/manual.shtml

Question: what does bowtie2-build do?

Picard's documentation also exists! Two bioinformatics programs in a row with decent documentation! Take a moment to celebrate, then have a look here: http://sourceforge.net/apps/mediawiki/picard/index.php

Question: Why not just remove all identitical pairs instead of mapping them and then removing them?

Question: What is the difference between samtools rmdup and Picard MarkDuplicates?

8.3.2 Mapping reads with bowtie2

First set up the files needed for mapping:

```
mkdir -p ~/mg-workshop/results/functional_annotation/mapping/$SAMPLE/
cd ~/mg-workshop/results/functional_annotation/mapping/$SAMPLE/
ln -s ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.1.fastq pair1.fastq
ln -s ~/mg-workshop/data/$SAMPLE/reads/1M/${SAMPLE_ID}_1M.2.fastq pair2.fastq
ln -s /proj/g2014180/nobackup/metagenomics-workshop/results/assembly/$SAMPLE/${SAMPLE}_31/contigs.fa
```

Then run the script that performs the mapping:

```
map-bowtie2-markduplicates.sh -t 16 -c pair1.fastq pair2.fastq $SAMPLE contigs.fa all map > map.log 2
```

8.3.3 Calculating coverage

We have now mapped reads back to the assembly and have information on how much of the assembly that is covered by the reads in the sample. What we are interested in is the coverage of the genes annotated in the previous steps by the PROKKA pipeline. To extract this information from the BAM file we first need to define the regions to calculate coverage for. This we will do by creating a custom GFF file defining the regions of interest (the PROKKA genes):

```
make.mapping.gff.sh ~/mg-workshop/results/functional_annotation/prokka/$SAMPLE/PROKKA_11252014.gff >
```

Next we extract coverage information from the BAM file for each gene in the GFF file we just created. We will use the bedtools coverage command within the BEDTools suite (https://code.google.com/p/bedtools/) that can parse a SAM/BAM file and a gff file to extract coverage information for every gene:

```
bedtools coverage -hist -abam map/all_$SAMPLE-smds.bam -b $SAMPLE.map.gff | sed "s/^M//g" > $SAMPLE.m
```

Have a look at the output file with less again. The final four columns give you the histogram i.e. coverage, number of bases with that coverage, length of the contig/feature/gene, bases with that coverage expressed as a ratio of the length of the contig/feature/gene.

To summarize the coverage for each gene we will use a script that calculates coverage from the histogram file you just produced:

```
get_coverage_for_genes.py -i <(echo $SAMPLE.map.hist) > $SAMPLE.coverage
```

We now have coverage values for all genes predicted and annotated by the PROKKA pipeline. Next, we will use the annotations and coverage values to summarize annotations for the sample. But first we will produce a file containing information on gene lengths in the sample. This will be used to normalize the gene coverage to make sure longer genes don't have higher coverage values just because of their length. Use the fastalength.py script for this:

```
fastalength.py -r -i ~/mg-workshop/results/functional_annotation/prokka/$SAMPLE/PROKKA_11252014.ffn >
```

Now we can move on to summarizing the annotations and producing interactive plots.

8.4 Summarize and explore the functional annotation

Now that we have annotated genes and quantified them in the sample using read mapping we will summarize and explore the annotations. This we will do by producing interactive plots detailing the proportion of functional categories such as metabolic pathways and orthologous gene families.

8.4.1 KRONA interactive plots

KRONA is a tool that takes as input a table of abundance values and several hierarchical categories and produces HTML files that can be explored interactively. The enzyme annotations from PROKKA are particularly suited for this purpose because these annotations can be grouped into higher functional categories such as pathways (e.g. glycolysis) and pathway classes (e.g. energy metabolism) for enzymes. Similarly, COG annotations can be summed up into higher categories such as "Carbohydrate transport and metabolism" and "Metabolism".

First we will create a new directory for the krona output and link to the necessary files:

```
mkdir -p ~/mg-workshop/results/functional_annotation/krona/$SAMPLE
cd ~/mg-workshop/results/functional_annotation/krona/$SAMPLE/
ln -s ~/mg-workshop/results/functional_annotation/mapping/$SAMPLE/$SAMPLE.coverage
ln -s ~/mg-workshop/results/functional_annotation/prokka/$SAMPLE/PROKKA.$SAMPLE.ec
ln -s ~/mg-workshop/results/functional_annotation/prokka/$SAMPLE/PROKKA.$SAMPLE.cog
ln -s ~/mg-workshop/results/functional_annotation/prokka/$SAMPLE/PROKKA.$SAMPLE.genelengths
```

```
ln -s ~/mg-workshop/results/functional_annotation/minpath/$SAMPLE/PROKKA.$SAMPLE.kegg.minpath
ln -s ~/mg-workshop/results/functional_annotation/minpath/$SAMPLE/PROKKA.$SAMPLE.metacyc.minpath
```

Next, use the genes.to.kronaTable.py script to produce the tabular output needed for KRONA.

For Metacyc pathways (from enzymes, only considering pathways predicted by MinPath):

```
genes.to.kronaTable.py -i PROKKA.$SAMPLE.ec -m ~/mg-workshop/reference_db/metacyc/ec.to.pwy -H ~/mg-w
```

For KEGG pathways (from enzymes, only considering pathways predicted by MinPath):

```
genes.to.kronaTable.py -i PROKKA.$SAMPLE.ec -m ~/mg-workshop/reference_db/kegg/ec.to.pwy -H ~/mg-worl
```

For COG annotations:

```
genes.to.kronaTable.py -i PROKKA.$SAMPLE.cog -m ~/mg-workshop/reference_db/cog/cog.to.cat -H ~/mg-wo
```

Then use the ktImportText script to generate the HTML files:

```
ktImportText -o $SAMPLE.krona.metacyc.minpath.html $SAMPLE.krona.metacyc.minpath.tab
ktImportText -o $SAMPLE.krona.kegg.minpath.html $SAMPLE.krona.kegg.minpath.tab
ktImportText -o $SAMPLE.krona.COG.html $SAMPLE.krona.COG.tab
```

Copy the resulting html files to your local computer with scp and open it a browser, like you did for the FastQC output.

Question: What are the main differences between the databases you have worked with: COG, Metacyc and KEGG? Which one do you prefer and why?

Question: What are the main differences between the different samples (gut, skin and teeth)? Compare with results from other groups. Can you for instance find differences in degradation of compounds?

Enjoy!